

GCP Professional Cloud Developer

Complete Study Guide

Professional Certification Guide

- ✓ Complete exam coverage
- ✓ Structured learning path
- ✓ Real-world scenarios
- ✓ Quick reference tables
- ✓ Best practices & tips

Version 1.0

Generated: December 22, 2025
© 2025 Study Materials

Table of Contents

GCP Professional Cloud Developer Study Guide

- Exam Snapshot
- Domain 1 Design scalable, available, and reliable applications (High)
 - Compute Service Decision Matrix
 - Real Exam Scenarios
 - Pub/Sub Patterns Deep Dive

Idempotent handler pattern

- Database Selection Matrix
- Common Exam Traps
- Domain 2 Build and test applications (High)
 - Testing Pyramid Strategy
 - Optimized Cloud Build Pipeline

cloudbuild.yaml - Production-ready example

- Step 1: Restore cache (speeds up npm/pip installs)**
- Step 2: Install dependencies**
- Step 3: Run linting (parallel with tests)**
- Step 4: Run unit tests (parallel)**
- Step 5: Run security scan (SAST)**
- Step 6: Build Docker image (wait for all tests)**
- Step 7: Scan image for vulnerabilities**
- Step 8: Push image**
- Step 9: Create provenance attestation (supply chain security)**
- Step 10: Update cache for next build**

- Secret Management Best Practices
- Vulnerability Scanning Matrix
- Private Pool Use Cases

Create private pool in VPC

Use in cloudbuild.yaml

- Build Performance Optimization
- Domain 3 Deploy applications (High)
- Domain 4 Integrate applications with Google Cloud services (MediumHigh)
- Domain 5 Secure applications (High)

- Domain 6 Monitor, log, and trace (High)
- High-Frequency Topic Map (What appears most)
- 30-Day Study Plan (2 hrs/day)
- Commands & Snippets

then split

- Cost Optimization Examples (Exam Favorite)
 - Scenario-Based Cost Analysis
 - Common Anti-Patterns (What NOT to Do)
 - Real-World Troubleshooting Scenarios
 - Problem 1: Cloud Run 503 errors under load

Check current concurrency setting

Reduce concurrency for single-threaded apps

OR make app multi-threaded

- Problem 2: Cloud Build timeout after 10 minutes

Option 1: Increase timeout

Option 2: Parallel steps (better)

Option 3: Caching (best)

Cache node_modules between builds

- Problem 3: Pub/Sub messages processed multiple times
- Problem 4: Cloud SQL connection exhaustion

Use Cloud SQL Python Connector with connection pooling

Limit Cloud Run instances

15 instances 5 conn = 75 connections (< 100 limit)

- Exam Strategy & Tips
 - Reading Questions
 - Decision Framework
 - Time Management
 - Red Flags (Wrong Answers)

GCP Professional Cloud Developer — Study Guide

Purpose: A structured, topic-by-topic plan aligned to the official exam domains, with focus indicators, hands-on tasks, and quick checks.



Exam Snapshot

- Format: 50–60 questions, 2 hours
 - Focus: Design, build, deploy, integrate, secure, and observe applications on GCP
 - Core services: Cloud Run, App Engine, GKE, Cloud Functions, Cloud Build, Artifact Registry, Pub/Sub, Cloud SQL/Firestore/Spanner, Memorystore, Cloud Logging/Monitoring, Secret Manager, IAM
-



Domain 1 — Design scalable, available, and reliable applications (High)

Topics and objectives:

- Workload placement: Cloud Run vs App Engine vs GKE vs Cloud Functions (latency, scaling, portability)
- Stateless vs stateful design; 12-factor app principles
- Resiliency patterns: retries, timeouts, circuit breakers, bulkheads
- Multi-region and regional HA patterns; load balancing options (Global HTTP(S) LB, Cloud Run domain mapping)
- Data design: choosing Cloud SQL vs Firestore vs Spanner; RTO/RPO; multi-region configs
- Event-driven architectures (Pub/Sub): at-least-once delivery, idempotency, ordering keys

■ Compute Service Decision Matrix

Requirement	Cloud Run	App Engine	GKE	Functions	
HTTP requests	Best	Good	Good	Limited	
Background tasks	Jobs	Tasks	CronJobs	Best	
Scale to zero	Yes	No (Std)	No	Yes	
Cold start	~1-3s	~10-30s	N/A (warm)	~1-5s	
Max timeout	60 min	60 min	Unlimited	9 min (2G)	
Stateful workloads	No	Limited	StatefulSet	No	
Custom runtime	Container	Limited	Any	Runtimes	
Pricing model	Request-based	Instance-hr	Node-hr	Invocation	
Multi-region	Manual	Built-in	Manual	Manual	
WebSocket support	Yes	Yes	Yes	No	
Streaming response	Yes	Yes	Yes	No	

■ Real Exam Scenarios

Scenario 1: "You need to deploy a containerized API that handles 10K requests/day with 95% coming during business hours. Cost is a primary concern."

- **Answer:** Cloud Run (scales to zero at night, pay-per-request)
- **Why not Functions:** Container requirement
- **Why not App Engine:** Can't scale to zero (standard)
- **Why not GKE:** Overkill, always-on costs

Scenario 2: "Legacy Java app needs migration with minimal changes. Requires background task queues and scheduled jobs."

- **Answer:** App Engine Standard (built-in Task Queues, Cron)
- **Why not Cloud Run:** Need to refactor for task queues
- **Why not GKE:** Over-engineering for simple app

Scenario 3: "Microservices with service mesh, Istio, and complex networking."

- **Answer:** GKE (full Kubernetes control)
- **Why not Cloud Run:** Limited networking control

Scenario 4: "Process uploaded images: resize, watermark, save to bucket."

- **Answer:** Cloud Functions (event-driven, triggered by Cloud Storage)

- **Why not Cloud Run:** Overkill for simple transformation

■ Pub/Sub Patterns Deep Dive

At-Least-Once Delivery:

- Messages may be delivered multiple times
- **Must implement idempotency:**
- Use unique message IDs to track processed messages
- Store message IDs in Firestore/Memorystore with TTL
- Example: payment processing (don't charge twice!)

```
# Idempotent handler pattern
import redis
redis_client = redis.Redis()

def process_message(message_id, data):
    # Check if already processed (TTL 7 days)
    if redis_client.exists(f"processed:{message_id}"):
        return {"status": "already_processed"}

    # Process the message
    result = do_business_logic(data)

    # Mark as processed
    redis_client.setex(f"processed:{message_id}", 604800, "1")
    return result
```

Ordering Keys:

- Use when message order matters for a specific entity
- Example: user profile updates (create → update → delete)
- All messages with same ordering key go to same partition

```
gcloud pubsub topics create orders --message-ordering
gcloud pubsub subscriptions create orders-sub \
--topic=orders --enable-message-ordering
```

Dead Letter Queues (DLQ):

- Poison messages that repeatedly fail processing
- After max attempts (5-100), move to DLQ for manual inspection
- Set up alerts when DLQ receives messages

■■ Database Selection Matrix

■ Use Case	■ Cloud SQL	■ Firestore	■ Spanner	■ Memorystore	■
■ Relational/SQL	■ ■ Postgres	■ ■ NoSQL	■ ■ Yes	■ ■ Cache	■
■ Global scale	■ ■ Regional	■ ■ Multi-reg	■ ■ Global	■ ■ Regional	■
■ Strong consist.	■ ■ Yes	■ ■ Optional	■ ■ Yes	■ ■ Yes	■
■ Transactions	■ ■ Full	■ ■ Limited	■ ■ Full	■ ■ Limited	■
■ Max DB size	■ 64 TB	■ Unlimited	■ Unlimited	■ 300 GB	■
■ Read replicas	■ ■ Yes	■ N/A	■ N/A	■ ■ Yes	■
■ Automatic shard	■ ■ No	■ ■ Yes	■ ■ Yes	■ ■ No	■
■ Cost (relative)	■ \$ Low	■ \$\$ Medium	■ \$\$\$ High	■ \$ Low	■
■ Best for	■ Legacy apps	■ Mobile/web	■ Finance/ERP	■ Session/cache	■

Key Decision Factors:

- **Cloud SQL:** Existing SQL app, < 10TB, single region, need read replicas
- **Firestore:** Mobile/web apps, flexible schema, multi-region reads, document model
- **Spanner:** Financial transactions, global consistency, unlimited scale, > 2TB
- **Memorystore:** Caching, sessions, pub/sub (Redis), sub-millisecond latency

■■ Common Exam Traps

1. **Trap:** "Use GKE for everything because it's most flexible"
 - **Reality:** Over-engineering, higher cost, more complexity
 - **Look for:** Simplest solution that meets requirements
2. **Trap:** "Firestore for all NoSQL needs"
 - **Reality:** Consider Bigtable for time-series, analytics workloads
3. **Trap:** "Always use multi-region for HA"
 - **Reality:** Regional can be sufficient (3 zones), lower latency/cost

Hands-on:

- Deploy a sample container to Cloud Run with min instances=1, concurrency=80, CPU throttling
- Set up Pub/Sub → Cloud Run push subscription with idempotent handler (Redis tracking)
- Create multi-region Cloud SQL instance with read replicas; test failover
- Implement circuit breaker pattern (exponential backoff, max retries)

Quick check:

- When to choose Spanner vs Cloud SQL? (Global scale vs regional, cost trade-off)

- How to guarantee idempotency with Pub/Sub? (Message ID tracking, Redis/Firestore)
 - What's the difference between Cloud Run min-instances and App Engine min-instances? (Billing model)
-

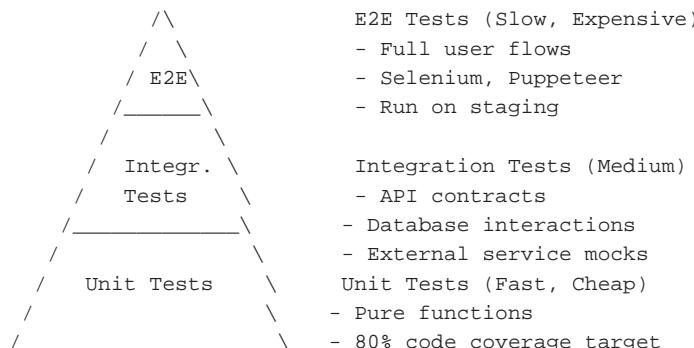


Domain 2 — Build and test applications (High)

Topics and objectives:

- Build systems: Cloud Build triggers, build steps, private pools, caching; `cloudbuild.yaml`
- Artifact management: Artifact Registry (Docker, Maven, npm), vulnerability scanning
- Testing strategies: unit/integration/e2e, contract tests, mocks; Cloud Build + test reports
- Local dev: Cloud Code, Skaffold, Cloud Run emulator, Functions Framework
- Secure coding: input validation, least privilege, service accounts per service, Secrets management

■ Testing Pyramid Strategy



■■ Optimized Cloud Build Pipeline

```
# cloudbuild.yaml - Production-ready example
options:
  machineType: 'E2_HIGHCPU_8' # Faster builds
  substitutionOption: 'ALLOW_LOOSE'
  logging: CLOUD_LOGGING_ONLY

substitutions:
  _REGION: 'us-central1'
  _ARTIFACT_REGISTRY: '${_REGION}-docker.pkg.dev/${PROJECT_ID}/app-repo'

steps:
# Step 1: Restore cache (speeds up npm/pip installs)
- name: 'gcr.io/cloud-builders/gsutil'
  args: ['cp', 'gs://${PROJECT_ID}_cloudbuild/cache/node_modules.tar.gz', '.']
  id: 'restore-cache'

# Step 2: Install dependencies
- name: 'node:18'
  entrypoint: 'bash'
  args:
    - '-c'
    - |
      if [ -f node_modules.tar.gz ]; then
        tar -xzf node_modules.tar.gz
      fi
      npm ci
  id: 'install-deps'
  waitFor: ['restore-cache']

# Step 3: Run linting (parallel with tests)
- name: 'node:18'
  entrypoint: 'npm'
  args: ['run', 'lint']
  id: 'lint'
  waitFor: ['install-deps']

# Step 4: Run unit tests (parallel)
- name: 'node:18'
  entrypoint: 'npm'
  args: ['run', 'test:unit']
  env:
    - 'CI=true'
  id: 'unit-tests'
  waitFor: ['install-deps']

# Step 5: Run security scan (SAST)
- name: 'aquasec/trivy'
  args: ['fs', '--severity', 'CRITICAL,HIGH', '--exit-code', '1', '.']
  id: 'security-scan'
  waitFor: ['install-deps']

# Step 6: Build Docker image (wait for all tests)
- name: 'gcr.io/cloud-builders/docker'
  args:
    - 'build'
    - '-t'
    - '${_ARTIFACT_REGISTRY}/app:${SHORT_SHA}'
    - '-t'
    - '${_ARTIFACT_REGISTRY}/app:latest'
    - '--cache-from'
    - '${_ARTIFACT_REGISTRY}/app:latest'
    - '.'
  id: 'build-image'
  waitFor: ['lint', 'unit-tests', 'security-scan']
```

```

# Step 7: Scan image for vulnerabilities
- name: 'gcr.io/cloud-builders/gcloud'
  entrypoint: 'bash'
  args:
  - '-c'
  - |
    gcloud artifacts docker images scan ${_ARTIFACT_REGISTRY}/app:${SHORT_SHA} \
      --location=${_REGION} --format=json
  id: 'image-scan'
  waitFor: ['build-image']

# Step 8: Push image
- name: 'gcr.io/cloud-builders/docker'
  args: ['push', '--all-tags', '${_ARTIFACT_REGISTRY}/app']
  id: 'push-image'
  waitFor: ['image-scan']

# Step 9: Create provenance attestation (supply chain security)
- name: 'gcr.io/cloud-builders/gcloud'
  entrypoint: 'bash'
  args:
  - '-c'
  - |
    gcloud beta container binauthz attestations sign-and-create \
      --artifact-url=${_ARTIFACT_REGISTRY}/app:${SHORT_SHA} \
      --attestor=prod-attestor \
      --attestor-project=${PROJECT_ID}
  id: 'attestation'
  waitFor: ['push-image']

# Step 10: Update cache for next build
- name: 'gcr.io/cloud-builders/gsutil'
  entrypoint: 'bash'
  args:
  - '-c'
  - |
    tar -czf node_modules.tar.gz node_modules
    gsutil cp node_modules.tar.gz gs://${PROJECT_ID}_cloudbuild/cache/
  id: 'save-cache'
  waitFor: ['push-image']

images:
- '${_ARTIFACT_REGISTRY}/app:${SHORT_SHA}'
- '${_ARTIFACT_REGISTRY}/app:latest'

timeout: '1200s' # 20 minutes

```

■ Secret Management Best Practices

■ Bad (Never do this):

```

env:
- DB_PASSWORD=mypassword123 # Exposed in logs/configs

```

■ Good (Secret Manager + Workload Identity):

```

steps:
- name: 'gcr.io/cloud-builders/gcloud'
  entrypoint: 'bash'
  args:
  - '-c'
  - |
    DB_PASSWORD=$(gcloud secrets versions access latest --secret=db-password)
    # Use $DB_PASSWORD in build steps
  secretEnv: ['DB_PASSWORD']

availableSecrets:
  secretManager:
  - versionName: projects/${PROJECT_ID}/secrets/db-password/versions/latest
    env: 'DB_PASSWORD'

```

Cloud Run with Secret Manager:

```

gcloud run deploy app \
--image=${IMAGE} \
--update-secrets=DB_PASSWORD=db-password:latest \
--service-account=app-sa@${PROJECT_ID}.iam.gserviceaccount.com

```

■ Vulnerability Scanning Matrix

Severity	Action	Who Gets Notified
CRITICAL	Block deployment	Security + Dev lead
HIGH	Review required	Dev lead
MEDIUM	Allow + ticket	Dev team
LOW	Allow (log)	Automated report

■ Private Pool Use Cases

1. **Access to VPC resources** (Cloud SQL, internal APIs)
2. **Larger machines** (custom machine types)
3. **Faster builds** (no cold start, persistent cache)
4. **Compliance** (data residency, private network)

```
# Create private pool in VPC
gcloud builds worker-pools create my-pool \
--region=us-central1 \
--peered-network=projects/${PROJECT_ID}/global/networks/default \
--worker-machine-type=e2-standard-4

# Use in cloudbuild.yaml
options:
  pool:
    name: 'projects/${PROJECT_ID}/locations/us-central1/workerPools/my-pool'
```

■ Build Performance Optimization

Technique	Time Saved	Complexity
Layer caching	30-50%	Low
Multi-stage builds	20-40%	Low
Parallel steps	20-60%	Medium
Private pool	10-20%	Medium
Dependency caching	40-70%	Medium
Larger machine type	30-50%	Low

Hands-on:

- Create optimized `cloudbuild.yaml` with parallel test execution and caching
- Enable Artifact Registry vulnerability scanning with severity policies
- Configure Binary Authorization with attestation requirement
- Use Secret Manager with per-environment secrets (dev/staging/prod)
- Set up private worker pool with VPC access to Cloud SQL

Quick check:

- GitHub App vs Cloud Source Repositories? (GitHub App = full CI/CD, CSR = simple Git)
- Purpose of private pool? (VPC access, custom machine, compliance)
- How to pass secrets safely? (Secret Manager, not env variables)
- What's the difference between `waitFor: ['-']` and omitting it? (Parallel vs sequential)



Domain 3 — Deploy applications (High)

Topics and objectives:

- Progressive delivery: blue/green, canary, traffic splitting (Cloud Run, App Engine), rollbacks
- Infra as Code: Terraform basics for app infra; Cloud Deploy for release orchestration
- Config and environment management: Config Controller, ConfigMaps/Secrets (K8s), service-to-service auth
- Zero-downtime migrations and DB schema versioning (Liquibase/Flyway)

Hands-on:

- Deploy two revisions in Cloud Run and split traffic (90/10 → 50/50 → 100/0)
- Define a Cloud Deploy pipeline to promote artifacts from dev → staging → prod
- Practice DB migration with a reversible change and rollback plan

Quick check:

- How would you implement a canary in Cloud Run and auto-rollback on errors?



Domain 4 — Integrate applications with Google Cloud services (Medium–High)

Topics and objectives:

- API exposure: API Gateway vs Cloud Endpoints; auth with JWTs/OIDC; rate limiting
- Messaging & streaming: Pub/Sub features (DLQs, ordering), Dataflow templates (overview)
- Data services: Cloud SQL connectivity (private IP/Serverless VPC Access), Firestore indexes, Spanner schemas
- Storage: Cloud Storage signed URLs, object lifecycle, event notifications
- Identity and access: IAM roles, custom roles, Workload Identity Federation, per-service identities

Hands-on:

- Protect a Cloud Run service behind API Gateway with JWT/Authenticator
- Use Serverless VPC Access to reach Cloud SQL from Cloud Run
- Configure Pub/Sub DLQ and test poison message handling

Quick check:

- When to choose API Gateway vs Endpoints?
 - How to set least-privilege for a build/deploy pipeline?
-



Domain 5 — Secure applications (High)

Topics and objectives:

- IAM best practices: least privilege, per-service accounts, minimal token scope, keyless auth
- Secret Manager vs env variables; KMS basics; CMEK for data services
- Organization policy constraints; Binary Authorization (GKE), Cloud Run invoker roles
- Supply chain security: Artifact signing (Sigstore), SBOMs, scanning

Hands-on:

- Replace env-based secrets with Secret Manager + per-revision access
- Enable Artifact Registry vulnerability scanning; block deploy on critical findings
- Configure Binary Authorization policy for GKE (if applicable)

Quick check:

- How do you rotate secrets without downtime?
-



Domain 6 — Monitor, log, and trace (High)

Topics and objectives:

- Observability stack: Cloud Logging, Cloud Monitoring, Error Reporting, Trace, Profiler, Debugger
- SLI/SLO/SLA basics; alerting policies; uptime checks; log-based metrics
- Distributed tracing for microservices; correlation IDs; structured logging
- Cost monitoring: quotas, budgets, labels

Hands-on:

- Emit structured JSON logs; create log-based metrics and alerts on error ratio
- Add OpenTelemetry auto-instrumentation to a service and view traces

- Define SLO (availability, latency) and attach alerting burn-rate policy

Quick check:

- What's a burn-rate alert and why is it used?
-



High-Frequency Topic Map (What appears most)

- Cloud Run vs GKE vs App Engine decision making — High
 - Cloud Build, Artifact Registry, CI test automation — High
 - Traffic splitting, canary/rollback, Cloud Deploy — High
 - Pub/Sub patterns: idempotency, DLQ, ordering — High
 - Cloud SQL/Firestore/Spanner selection and connectivity — High
 - IAM, Secret Manager, secure service-to-service auth — High
 - Observability: logs, metrics, traces, SLOs, alerts — High
-



30-Day Study Plan (2 hrs/day)

- Week 1: Workload placement, resiliency, Pub/Sub patterns; Hands-on Cloud Run + Pub/Sub
 - Week 2: CI/CD with Cloud Build, Artifact Registry; Security with Secret Manager; Canary deploys
 - Week 3: Data services (SQL/Firestore/Spanner), VPC access, API Gateway; Observability stack
 - Week 4: End-to-end project, SLOs + burn-rate alerts; Review sample questions and weak areas
-



Commands & Snippets

- Cloud Build trigger:

```
steps:  
- name: 'gcr.io/cloud-builders/docker'  
  args: [ 'build', '-t', '$REGION-docker.pkg.dev/$PROJECT/repo/app:$COMMIT_SHA', '.' ]  
- name: 'gcr.io/cloud-builders/docker'  
  args: [ 'push', '$REGION-docker.pkg.dev/$PROJECT/repo/app:$COMMIT_SHA' ]  
images:  
- '$REGION-docker.pkg.dev/$PROJECT/repo/app:$COMMIT_SHA'
```

- Cloud Run deploy with traffic split:

```
gcloud run deploy app \  
--image=$REGION-docker.pkg.dev/$PROJECT/repo/app:$COMMIT_SHA \  
--region=$REGION --platform=managed --allow-unauthenticated  
# then split  
gcloud run services update-traffic app \  
--to-revisions REV1=90,REV2=10 --region=$REGION
```

- Pub/Sub DLQ policy example:

```
gcloud pubsub subscriptions create sub \  
--topic=topic --dead-letter-topic=dlq --max-delivery-attempts=5
```



■ Cost Optimization Examples (Exam Favorite)

Scenario-Based Cost Analysis

Scenario 1: API Service (10K requests/day, 50ms avg)

Cloud Run (scale-to-zero):

- Requests: 10K/day = 300K/month
- vCPU: 1 vCPU @ \$0.00002400/vCPU-sec
- Memory: 512MB @ \$0.00000250/GB-sec
- Request: \$0.40/million
- Monthly: ~\$8-12

App Engine Standard:

- Instance hours: ~720 hrs (always 1 instance)
- F1 instance: ~\$0.05/hr
- Monthly: ~\$36

GKE:

- Node: e2-medium (2vCPU, 4GB) = \$24.27/mo
- Cluster fee: \$73/mo
- Monthly: ~\$97

■ Winner: Cloud Run (saves 70-90%)

Scenario 2: Background Job (1 hour/day, CPU-intensive)

Cloud Run Jobs:

- 1 vCPU, 2GB, 1 hour/day
- Cost: ~\$3/month

GKE CronJob:

- Always-on node even if job runs 1hr
- Cost: ~\$97/month (node + cluster)

■ Winner: Cloud Run Jobs (saves 95%)

Scenario 3: Stateful Application (24/7, persistent storage)

GKE with StatefulSet:

- Right choice for stateful workloads
- Persistent volumes
- Cost: \$97/mo + storage

Cloud Run:

- ■ Cannot persist state between requests
- Wrong choice

■ Winner: GKE (only option)



■ Common Anti-Patterns (What NOT to Do)

1. **Anti-pattern:** Storing secrets in environment variables
 - **Problem:** Visible in logs, console, metadata
 - **Solution:** Secret Manager with IAM
 2. **Anti-pattern:** Using default service account
 - **Problem:** Over-privileged (Editor role)
 - **Solution:** Custom service account per service
 3. **Anti-pattern:** Polling Pub/Sub in a loop
 - **Problem:** Wastes CPU, increases latency
 - **Solution:** Push subscription to Cloud Run
 4. **Anti-pattern:** Building images on local machine
 - **Problem:** "Works on my machine", no security scanning
 - **Solution:** Cloud Build with automated scanning
 5. **Anti-pattern:** Hard-coding configuration
 - **Problem:** Different values per environment
 - **Solution:** Environment variables or Secret Manager
 6. **Anti-pattern:** No retry logic for external calls
 - **Problem:** Transient failures break application
 - **Solution:** Exponential backoff with max retries
-



■ Real-World Troubleshooting Scenarios

Problem 1: Cloud Run 503 errors under load

Symptoms:

- Works fine with low traffic
- 503 "Service Unavailable" at 100+ req/sec
- CPU not maxed out

Diagnosis:

```
# Check current concurrency setting
gcloud run services describe app --region=us-central1 \
--format='value(spec.template.spec.containerConcurrency)'
```

Root cause: Default concurrency=80, but your app is single-threaded

Solution:

```
# Reduce concurrency for single-threaded apps
gcloud run services update app --concurrency=1 --region=us-central1
# OR make app multi-threaded
```

Problem 2: Cloud Build timeout after 10 minutes

Symptoms:

- Build fails with "Build timeout"
- npm install takes 8 minutes
- Tests take 5 minutes

Diagnosis: Default timeout is 10 minutes

Solutions:

```
# Option 1: Increase timeout
timeout: '3600s' # 1 hour

# Option 2: Parallel steps (better)
steps:
- name: 'node'
  args: ['npm', 'install']
  id: 'install'
- name: 'node'
  args: ['npm', 'run', 'test:unit']
  waitFor: ['install'] # Parallel after install
  id: 'test-unit'
- name: 'node'
  args: ['npm', 'run', 'test:integration']
  waitFor: ['install'] # Parallel
  id: 'test-int'

# Option 3: Caching (best)
# Cache node_modules between builds
```

Problem 3: Pub/Sub messages processed multiple times

Symptoms:

- Payment charged twice
- Database records duplicated
- Message count shows delivery but backlog increases

Root cause: Not idempotent, acknowledgement failures

Solution:

```
import hashlib
from google.cloud import firestore

db = firestore.Client()

def process_message(message):
    message_id = message.message_id
    doc_ref = db.collection('processed').document(message_id)

    # Atomic check-and-set
    transaction = db.transaction()
    snapshot = doc_ref.get(transaction=transaction)

    if snapshot.exists:
        return {'status': 'duplicate', 'message_id': message_id}

    # Process message
    result = do_business_logic(message.data)

    # Mark as processed (atomic)
    doc_ref.set({'processed_at': firestore.SERVER_TIMESTAMP})

    return result
```

Problem 4: Cloud SQL connection exhaustion

Symptoms:

- "Too many connections" error
- Cloud Run scales up, database fails
- Max connections: 100

Root cause: Each Cloud Run instance creates connection pool

Solutions:

1. **Connection pooling:**

```
# Use Cloud SQL Python Connector with connection pooling
from google.cloud.sql.connector import Connector
import sqlalchemy

connector = Connector()

def getconn():
    return connector.connect(
        "project:region:instance",
        "pg8000",
        pool_size=5, # Limit per instance
        max_overflow=2
    )

pool = sqlalchemy.create_engine(
    "postgresql+pg8000://",
    creator=getconn,
    pool_size=5, # Max 5 connections per instance
    max_overflow=2
)
```

2. Set max-instances:

```
# Limit Cloud Run instances
gcloud run services update app \
    --max-instances=15 \
    --region=us-central1
# 15 instances x 5 conn = 75 connections (< 100 limit)
```

3. Use Cloud SQL Proxy:

- Handles connection pooling automatically
- Encrypted connections



■ Exam Strategy & Tips

Reading Questions

1. Highlight key constraints:

- Budget concerns → Prefer serverless
- Compliance/regulations → Organization policies, CMEK
- "Minimal changes" → Lift-and-shift (App Engine, GKE)
- "Cloud-native" → Cloud Run, Functions

- "Existing K8s" → GKE

2. Watch for keywords:

- "Cost-effective" = Cloud Run > App Engine > GKE
- "Low latency" = Set min-instances, use Memorystore
- "Secure" = Secret Manager, private endpoints, least privilege
- "Scalable" = Autoscaling, stateless, managed services
- "Reliable" = Multi-region, retry logic, error budgets

Decision Framework

Question: How to deploy X?

1. What's the constraint?
 - Cost → Serverless
 - Control → GKE
 - Simplicity → Cloud Run
 - Legacy → App Engine
2. What's the workload?
 - HTTP API → Cloud Run
 - Events → Functions
 - Batch → Cloud Run Jobs
 - Stateful → GKE
3. What's the scale?
 - Variable → Autoscaling (Run/Functions)
 - Steady → App Engine
 - Huge → GKE + HPA

Time Management

- 2 hours for 50-60 questions = ~2 minutes per question
- **Strategy:**
 - First pass: Answer easy questions (30-40 min)
 - Mark difficult ones for review
 - Second pass: Tackle marked questions (40-50 min)
 - Final review: Check marked answers (20-30 min)

Red Flags (Wrong Answers)

- ■ Over-engineering ("Use GKE for simple API")

- ■ Ignoring managed services ("Build your own load balancer")
 - ■ Security holes ("Store keys in code")
 - ■ No error handling ("Just deploy and hope")
 - ■ Ignoring cost ("Use biggest machines")
-



Final Review Checklist

- [] Clear on workload placement trade-offs
- [] Can build/test artifacts and enforce scanning
- [] Can deploy progressively and rollback safely
- [] Can integrate with managed data services securely
- [] Emits structured logs and traces; SLOs with alerts
- [] Familiar with sample questions and timing